

# Real-Time Face Blur

- ◆ Project 2: Real-Time Face Blur for Privacy

Usefulness: Used in CCTV feeds or classroom monitoring to anonymize faces.

Core CVIP concepts:

Haar Cascade detection (OpenCV built-in)

Region of Interest (ROI)

Gaussian blurring / pixelation

Steps:

Detect faces using OpenCV's pretrained Haar Cascade.

Extract ROI of each detected face.

Apply cv2.GaussianBlur() or pixelate using cv2.resize()  
down and up.

Replace the blurred region back into the frame.

Show live webcam feed with blurred faces.

✓ Meaningful because: it demonstrates how CVIP techniques can preserve privacy without ML or cloud APIs.

expalin everything

```
# Alternative version with more optimizations
import cv2
import time
```

```
# Load Haar Cascade
face_cascade =
```

```

cv2.CascadeClassifier(cv2.data.haarcascades +
'haarcascade_frontalface_default.xml')

# Webcam setup
cap = cv2.VideoCapture(0)
cap.set(cv2.CAP_PROP_FRAME_WIDTH, 640)
cap.set(cv2.CAP_PROP_FRAME_HEIGHT, 480)
cap.set(cv2.CAP_PROP_FPS, 30)

# Performance settings
frame_skip = 3
frame_counter = 0
prev_faces = []
use_prev_faces = True # Flag to use previous detection

# Region of interest for face detection (reduce search
area)
roi_enabled = False
roi_coords = (0, 0, 640, 480) # x, y, w, h

# Timing for performance monitoring
last_time = time.time()
fps = 0

while True:
    ret, frame = cap.read()
    if not ret:
        break

    current_time = time.time()
    fps = 0.9 * fps + 0.1 * (1 / (current_time - last_time)) if
    current_time - last_time > 0 else fps
    last_time = current_time

    frame_counter += 1

    # Face detection on selected frames
    if frame_counter % frame_skip == 0 or not
    use_prev_faces:
        # Convert to grayscale
        gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

```

```

# Apply ROI if enabled
if roi_enabled:
    x, y, w, h = roi_coords
    roi_gray = gray[y:y+h, x:x+w]
    faces = face_cascade.detectMultiScale(roi_gray,
scaleFactor=1.1, minNeighbors=5, minSize=(50, 50))
    # Adjust coordinates to full frame
    faces = [(x+fx, y+fy, fw, fh) for (fx, fy, fw, fh) in
faces]
else:
    faces = face_cascade.detectMultiScale(gray,
scaleFactor=1.1, minNeighbors=5, minSize=(50, 50))

prev_faces = faces
use_prev_faces = True
else:
    faces = prev_faces

# Process faces
for (x, y, w, h) in faces:
    # Extract and blur face
    face_roi = frame[y:y+h, x:x+w]
    blurred_face = cv2.GaussianBlur(face_roi, (45, 45),
30)
    frame[y:y+h, x:x+w] = blurred_face

    # Draw rectangle
    cv2.rectangle(frame, (x, y), (x+w, y+h), (0, 255, 0), 2)

    # Update ROI for next detection (face tracking
optimization)
    if roi_enabled:
        roi_coords = (max(0, x-50), max(0, y-50), min(640,
w+100), min(480, h+100))

    # Display FPS
    cv2.putText(frame, f"FPS: {fps:.1f}", (10, 30),
cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)

    cv2.imshow("Optimized Face Blur", frame)

key = cv2.waitKey(1) & 0xFF

```

```
if key == 27:  
    break  
elif key == ord('r'): # Reset ROI  
    roi_enabled = not roi_enabled  
    roi_coords = (0, 0, 640, 480)  
    use_prev_faces = False  
  
cap.release()  
cv2.destroyAllWindows()
```

Alright Kush — let's dissect this project **line by line and concept by concept**, so you know *exactly* what's happening and why this project is a solid, meaningful CVIP demonstration.

---

## ◆ Project Overview

**Name:** Real-Time Face Blur for Privacy

**Goal:** Detect faces in a live webcam feed and blur them in real-time — useful for **privacy-preserving video streams** (CCTV, classrooms, etc.).

**Core CVIP Concepts:**

- **Face Detection (Haar Cascade)**
  - **Region of Interest (ROI)**
  - **Gaussian Blurring / Pixelation**
  - **Frame Optimization Techniques**
- 

## ◆ Step-by-Step Explanation

### 1 Importing Libraries

```
python
```

```
import cv2  
import time
```

- **cv2** → Core OpenCV library for image processing.
  - **time** → Used for calculating FPS (frames per second) for performance measurement.
-

## 2 Load Pretrained Haar Cascade Classifier

python

```
face_cascade = cv2.CascadeClassifier(cv2.data.haarcascades +  
'haarcascade_frontalface_default.xml')
```

- OpenCV provides pre-trained **Haar Cascade** XML models for face, eyes, smiles, etc.
- `haarcascade_frontalface_default.xml` detects human faces.
- It's fast and works on grayscale images.

### 🧠 Concept:

Haar Cascade uses simple rectangular features + AdaBoost + sliding window to detect patterns similar to human faces.

---

## 3 Webcam Initialization

python

```
cap = cv2.VideoCapture(0)  
cap.set(cv2.CAP_PROP_FRAME_WIDTH, 640)  
cap.set(cv2.CAP_PROP_FRAME_HEIGHT, 480)  
cap.set(cv2.CAP_PROP_FPS, 30)
```

- Opens webcam (`0` = default camera).
- Sets resolution and frame rate.
- Balanced setup for **speed + clarity**.

## 4 Performance Optimization Variables

python

```
frame_skip = 3  
frame_counter = 0  
prev_faces = []  
use_prev_faces = True
```

- **frame\_skip**: Detect faces only every 3rd frame — reduces heavy computation.
- **prev\_faces**: Stores face coordinates from the last detection frame.

- **use\_prev\_faces**: Uses old detections temporarily to avoid lag.
- 

## 5 ROI Setup (Optional)

```
python
```

```
roi_enabled = False  
roi_coords = (0, 0, 640, 480)
```

- **ROI (Region of Interest)**: Restrict search area to part of the frame → speeds up detection if you're tracking a known area (like a classroom zone).

When enabled, detection only happens inside ROI boundaries.

---

## 6 FPS Tracking

```
python
```

```
last_time = time.time()  
fps = 0
```

Used to measure how many frames per second your system is processing, helpful for optimizing CVIP systems.

---

## 7 Main Loop

```
python
```

```
while True:  
    ret, frame = cap.read()  
    if not ret:  
        break
```

Reads one frame at a time from webcam.

If reading fails (camera disconnected), it exits.

---

## 8 FPS Calculation (Exponential Smoothing)

```
python
```

```
current_time = time.time()  
fps = 0.9 * fps + 0.1 * (1 / (current_time - last_time)) if current_time - last_time > 0 else fps  
last_time = current_time
```

- Calculates instantaneous FPS and smooths it using exponential average to avoid flickering values.
- 

## 9 Controlled Face Detection

```
python
```

```
if frame_counter % frame_skip == 0 or not use_prev_faces:  
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
```

- Converts frame to grayscale (Haar classifier works only on grayscale images).
  - Detects faces only every few frames to boost performance.
- 

## 🔍 Face Detection Process

```
python
```

```
faces = face_cascade.detectMultiScale(  
    gray,  
    scaleFactor=1.1,  
    minNeighbors=5,  
    minSize=(50, 50)  
)
```

Parameters:

- **scaleFactor=1.1**: How much the image size is reduced at each scale. (Lower = more accurate but slower.)
- **minNeighbors=5**: How many detections are required for it to be considered a valid face.
- **minSize**: Ignores tiny detections.

This returns a list of `(x, y, w, h)` coordinates for each detected face.

---

## 1 0 ROI Adjustment (if enabled)

If ROI was active, it detects within a smaller cropped zone, then shifts coordinates to full-frame.

---

## 1 1 Blurring Detected Faces

python

```
for (x, y, w, h) in faces:  
    face_roi = frame[y:y+h, x:x+w]  
    blurred_face = cv2.GaussianBlur(face_roi, (45, 45), 30)  
    frame[y:y+h, x:x+w] = blurred_face
```

- Extract the **face region (ROI)**.
- Apply **Gaussian blur**: uses a weighted kernel (here 45×45 pixels) to smooth and obscure details.
- Replace original face with blurred face.

🧠 Alternative:

You could **pixelate** instead by resizing down and back up:

python

```
small = cv2.resize(face_roi, (10, 10))  
frame[y:y+h, x:x+w] = cv2.resize(small, (w, h), interpolation=cv2.INTER_NEAREST)
```

---

## 1 2 Optional Visualization

python

```
cv2.rectangle(frame, (x, y), (x+w, y+h), (0, 255, 0), 2)
```

Draws rectangles for debugging or visualization.

---

## 1 3 Display FPS on Frame

python

```
cv2.putText(frame, f"FPS: {fps:.1f}", (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 0.7, (0, 255, 0), 2)
```

Shows current FPS in the corner — useful for testing optimization effects.

---

## 1 4 Display Output

```
python
```

```
cv2.imshow("Optimized Face Blur", frame)
```

Displays live video with blurred faces.

---

## 1 5 Keyboard Controls

```
python
```

```
key = cv2.waitKey(1) & 0xFF
if key == 27:
    break
elif key == ord('r'):
    roi_enabled = not roi_enabled
    roi_coords = (0, 0, 640, 480)
    use_prev_faces = False
```

- Press **Esc** → quit
  - Press **R** → toggle ROI tracking mode
- 

## 1 6 Clean Exit

```
python
```

```
cap.release()
cv2.destroyAllWindows()
```

Closes webcam and OpenCV windows cleanly.

---

- ◆ **Why This Project Is Meaningful**

1. **Privacy-preserving application** — directly relevant to CCTV and smart surveillance.
  2. **Demonstrates key CVIP pipeline:**
    - Capture → Detect → Process → Display
  3. **Efficient** — uses frame skipping, ROI tracking, and cached detections.
  4. **No heavy ML model** — uses OpenCV's classical methods.
  5. **Can be extended easily** — motion-based tracking, dynamic ROI, or adding person anonymization.
- 

- ◆ **Extensions / Improvements**

If you want to push this further:

- Replace Haar with **DNN-based face detector** (like `opencv_face_detector_uint8.pb`).
  - Add a toggle to switch between **blur and pixelation**.
  - Store blurred output video using `cv2.VideoWriter`.
  - Integrate motion detection → blur only moving faces.
  - Use **YOLOv8 face model** for better accuracy.
- 

Would you like me to explain the **mathematical concept** behind Gaussian blur (kernel convolution and standard deviation effect) next, or the **difference between Haar Cascade and modern DNN face detectors** — which would help you justify your CVIP approach in a report?